



Introduction of Reinforcement Learning

Artificial Intelligence

- 지능이란?

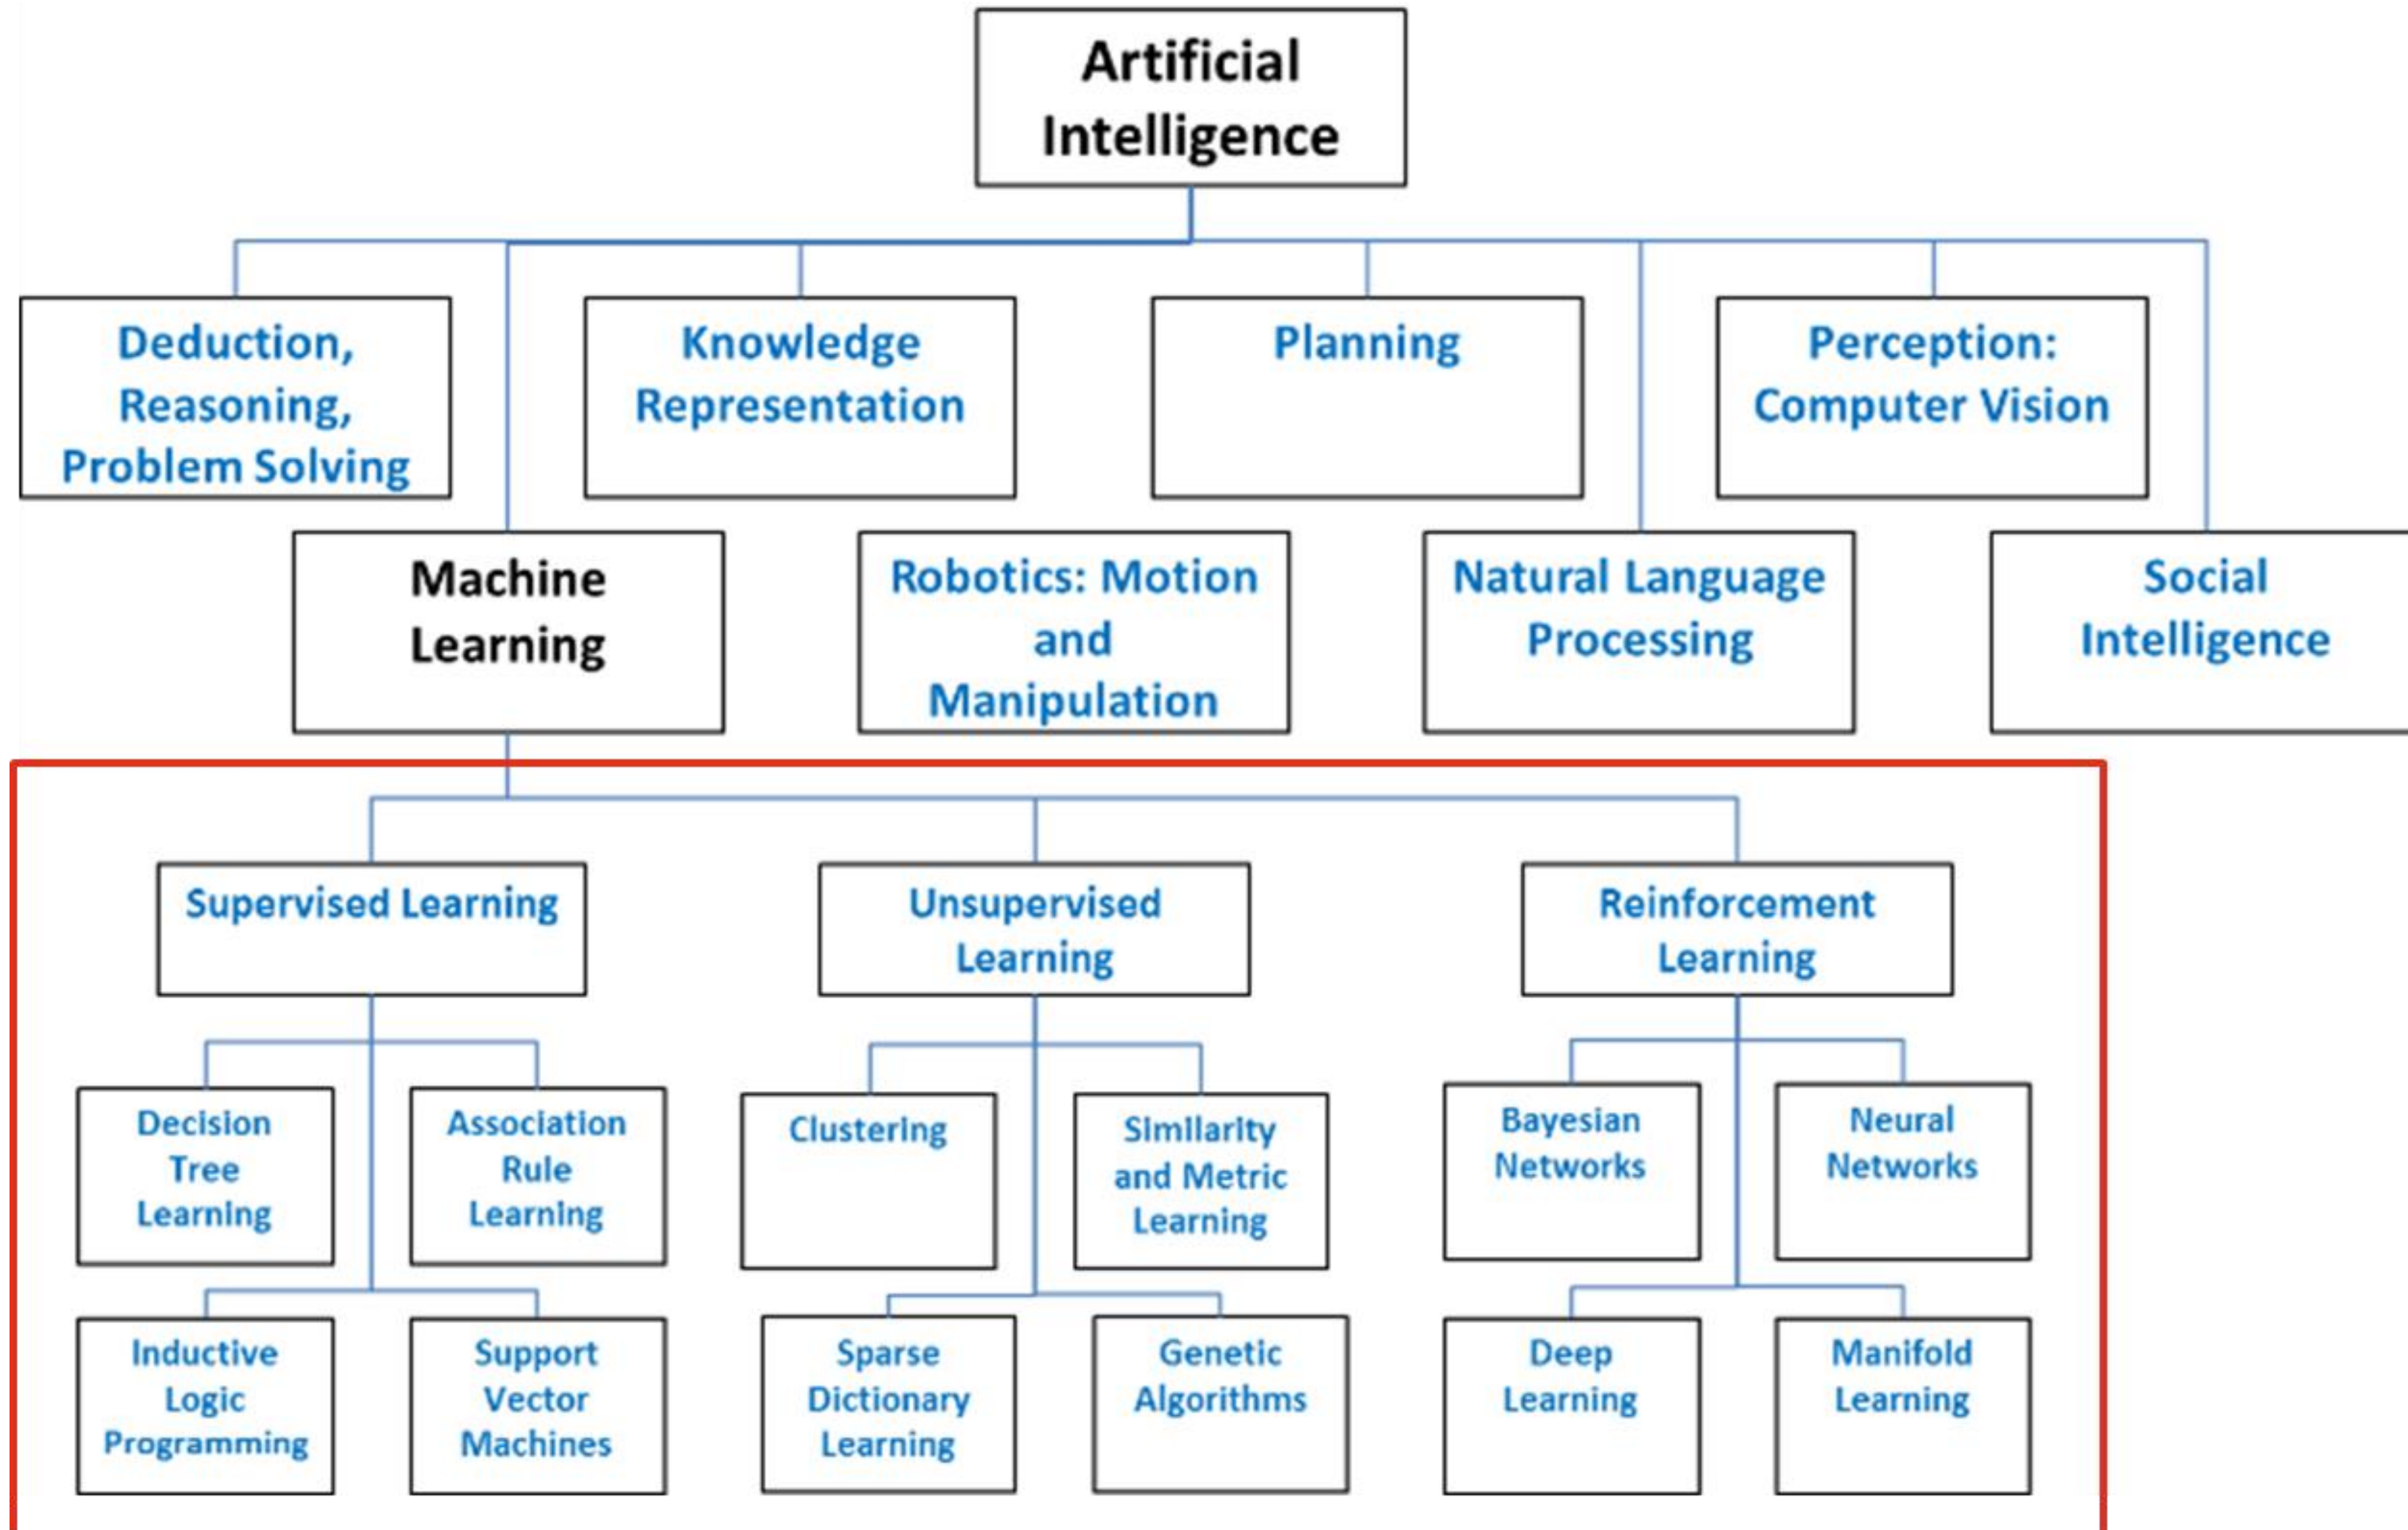
→ 보다 추상적인 정보를 이해하는 능력



- 인공 지능이란?

→ 이러한 지능 현상을 인공적으로 구현하려는 연구

Artificial Intelligence & Machine Learning




Deep Learning in RL




Discrete Choices

⋮



Layer 2 Features



Layer 1 Features



Original Data

- DeepRL에서 딥러닝은 그저 하나의 module로써만 사용된다.

- **Deep Learning의 강점 :**

1. 층을 쌓을 수록 더욱 Abstract Feature Learning이 가능한 유일한 알고리즘.
2. Universal Function Approximator.

Universal Function Approximator

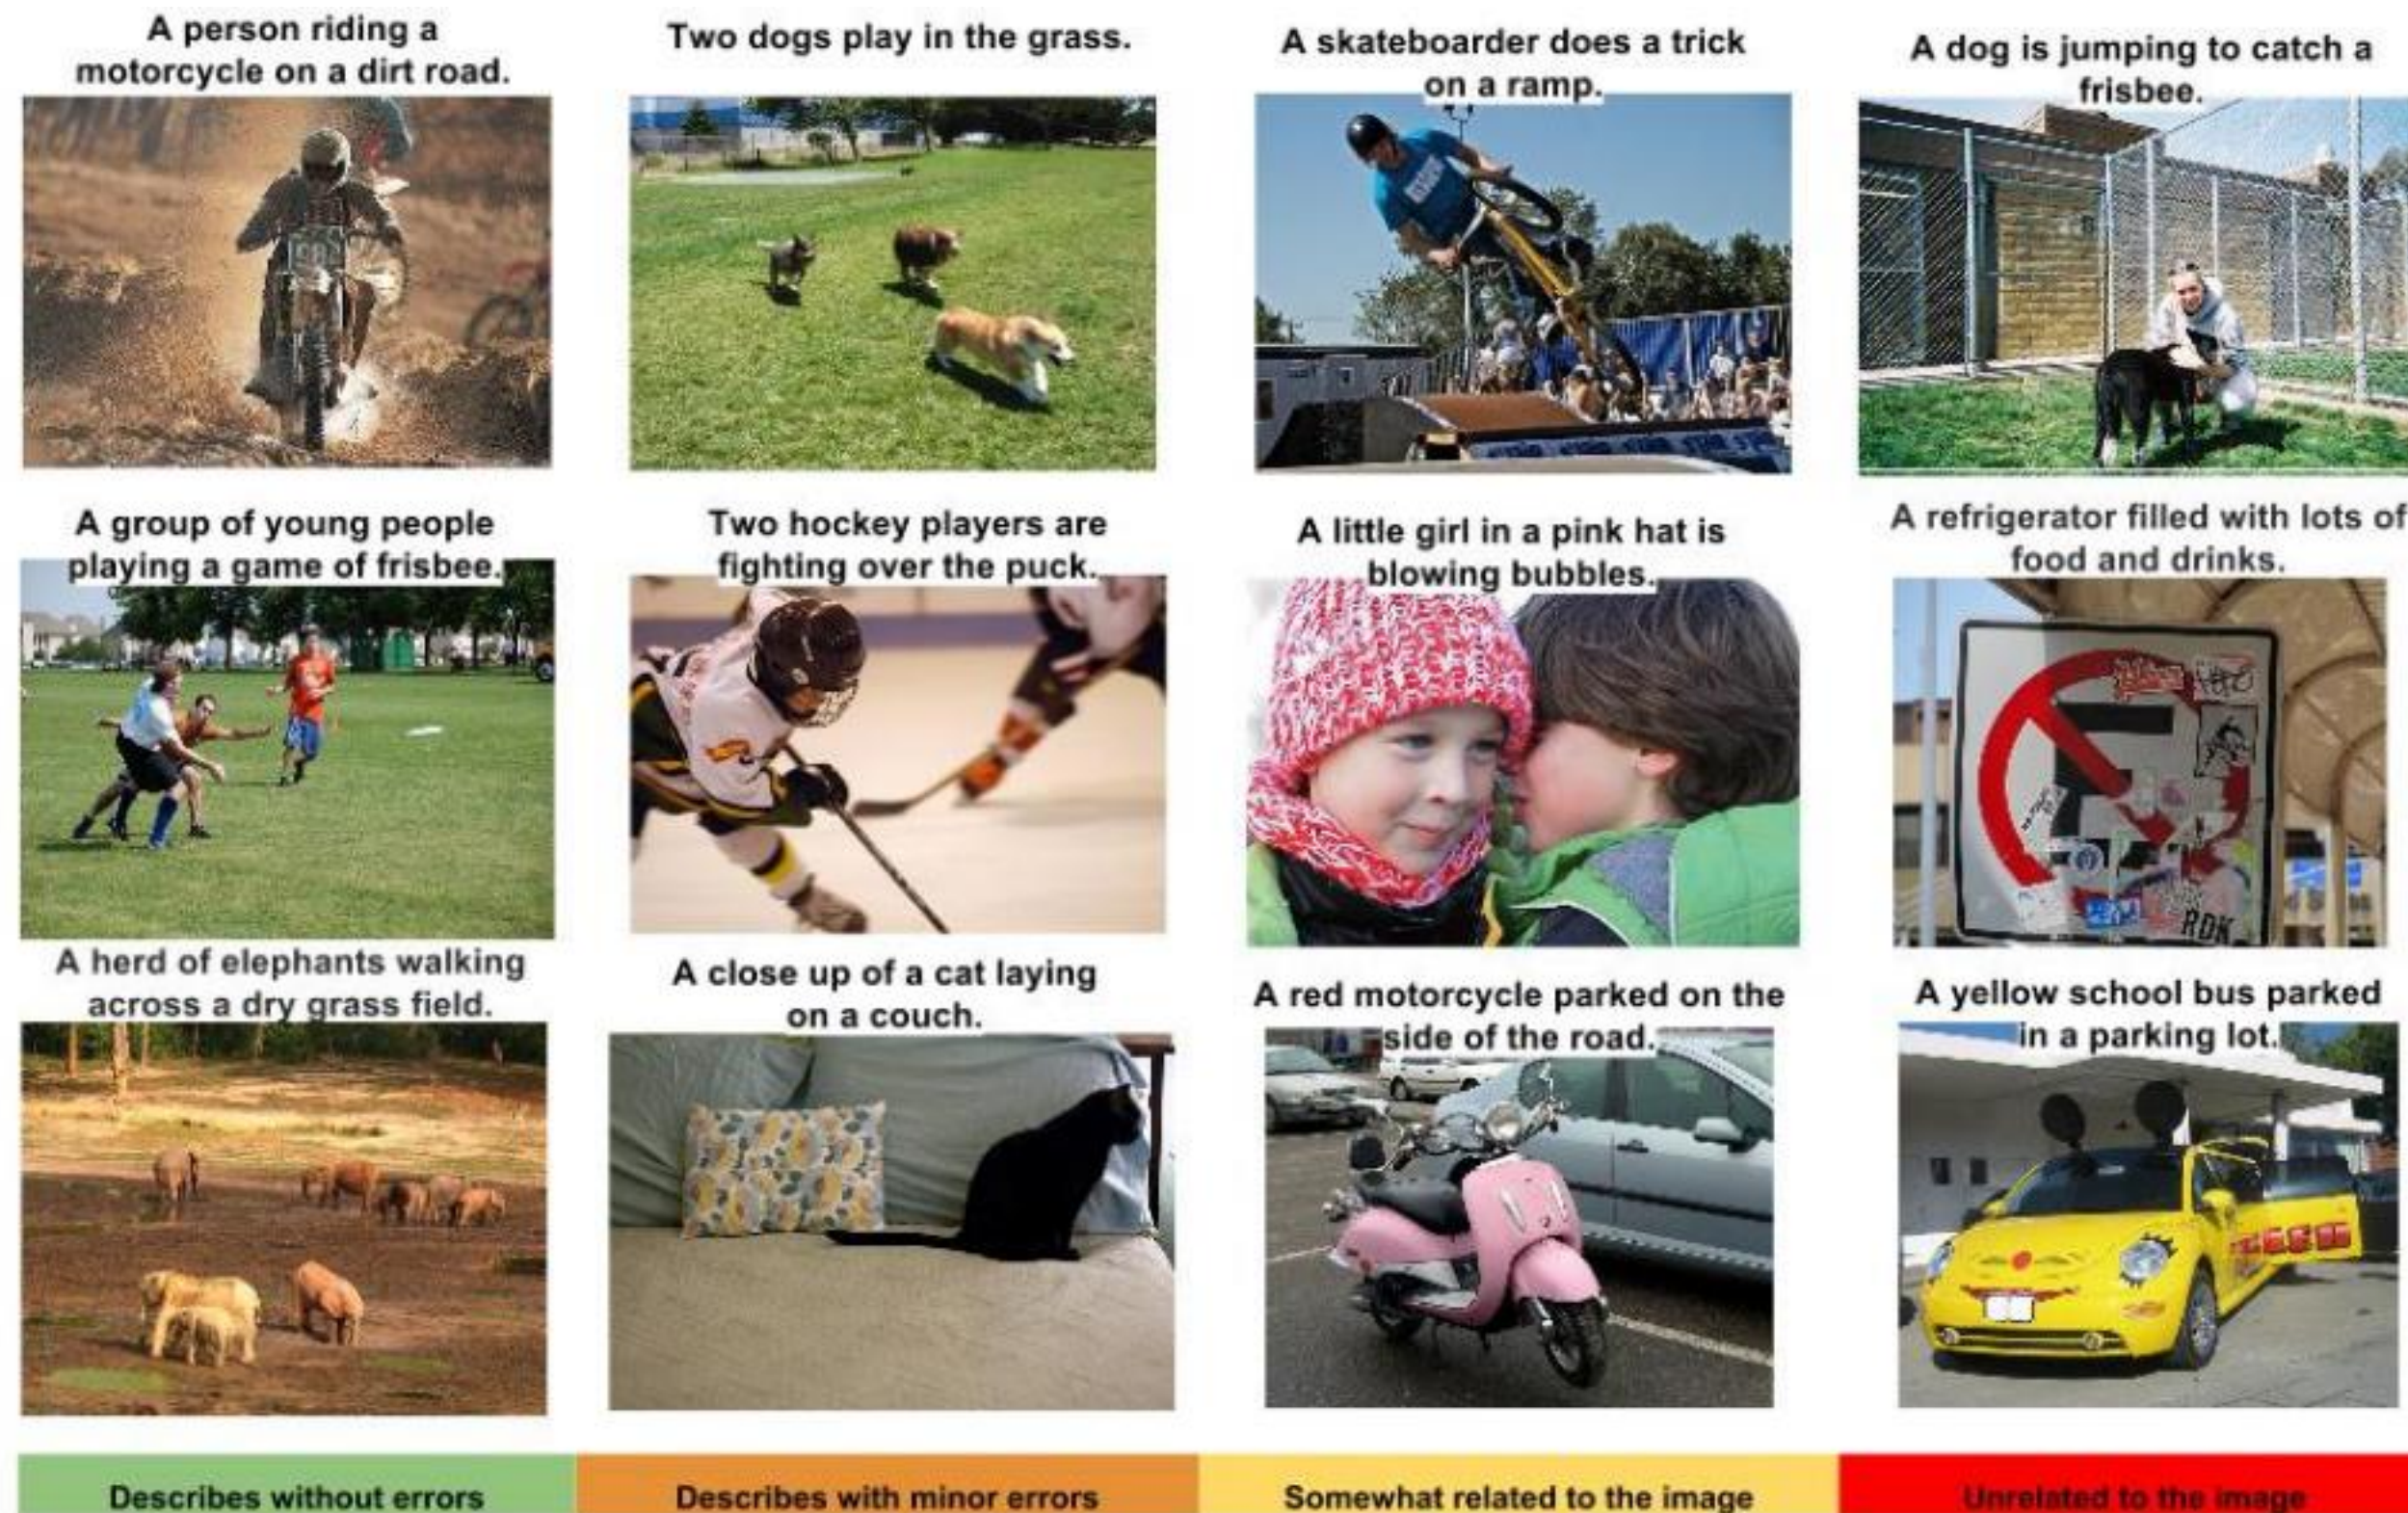



Figure 5. A selection of evaluation results, grouped by human rating.



Reinforcement Learning이란?

Machine Learning

- Supervised Learning :

$$y = f(x)$$


- Unsupervised Learning :

$$x \sim p(x) \quad \text{or} \quad x = f(x)$$

- Reinforcement Learning :

Find a policy, $p(a|s)$ which maximizes the sum of reward

Example of Supervised Learning : Polynomial Curve Fitting



추세선 서식

추세선 옵션

선 색

선 스타일

그림자

추세/회귀 유형

☐ 지수(X)

☒ 선형(L)

☐ 로그(Q)

☐ 다항식(P) 차수(D): 2

☐ 거듭제곱(W)

☐ 이동 평균(M) 구간(E): 2

추세선 이름

☒ 자동(A): 선형 (계열1)

☐ 사용자 지정(C):

예측

앞으로(F): 0,0 구간

뒤로(B): 0,0 구간

☐ 절편(S) = 0,0


☐ 수식을 차트에 표시(E)

☐ R-제곱 값을 차트에 표시(B)


닫기

Microsoft Excel 2007의 추세선


Example of Unsupervised Learning : **Clustering**




Example of Reinforcement Learning : Optimal Control Problem



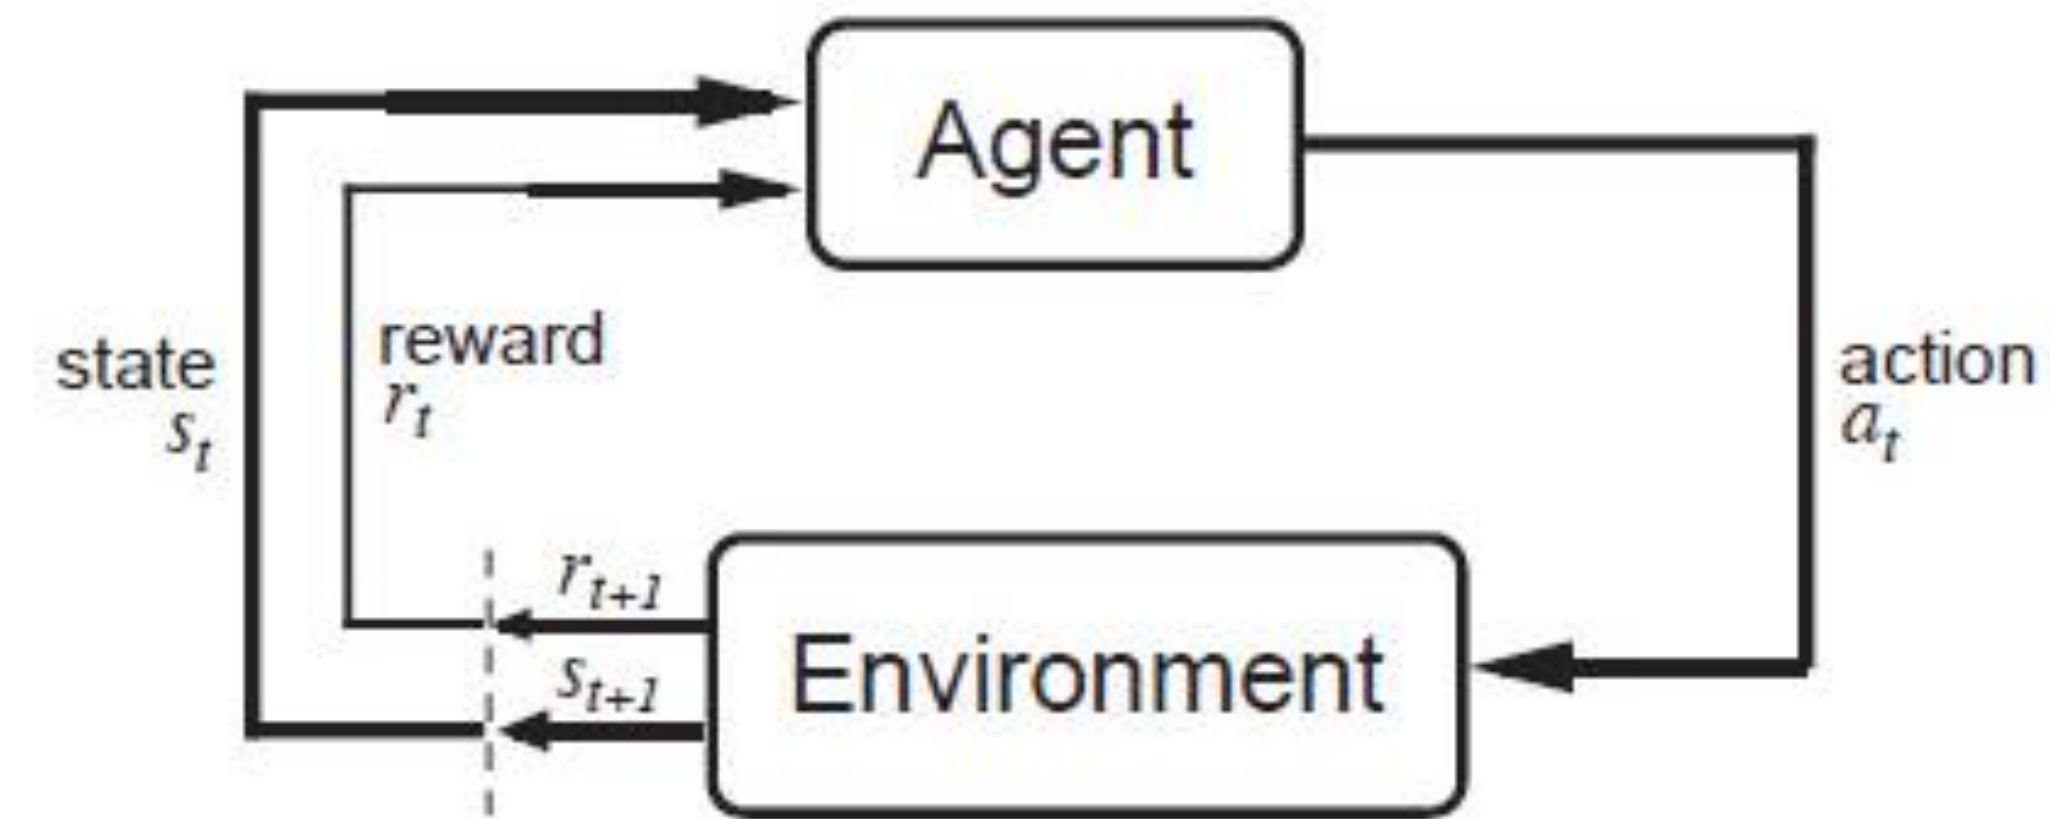

State 1



State 2




State 3



Markov Decision Processes(MDP)


Markov Processes

- Discrete state space : $S = \{ A , B \}$
- State transition probability : $P(S'|S) = \{ P_{AA} = 0.7 , P_{AB} = 0.3 , P_{BA} = 0.5 , P_{BB} = 0.5 \}$
- Purpose : Finding a steady state distribution




Markov Reward Processes

- Discrete state space : $S = \{ A , B \}$
- State transition probability : $P(S'|S) = \{ P_{AA} = 0.7 , P_{AB} = 0.3 , P_{BA} = 0.5 , P_{BB} = 0.5 \}$
- Reward function : $R(S'=A) = +1 , R(S'=B) = -1$
- Purpose : Finding an expected reward distribution



Markov Decision Processes

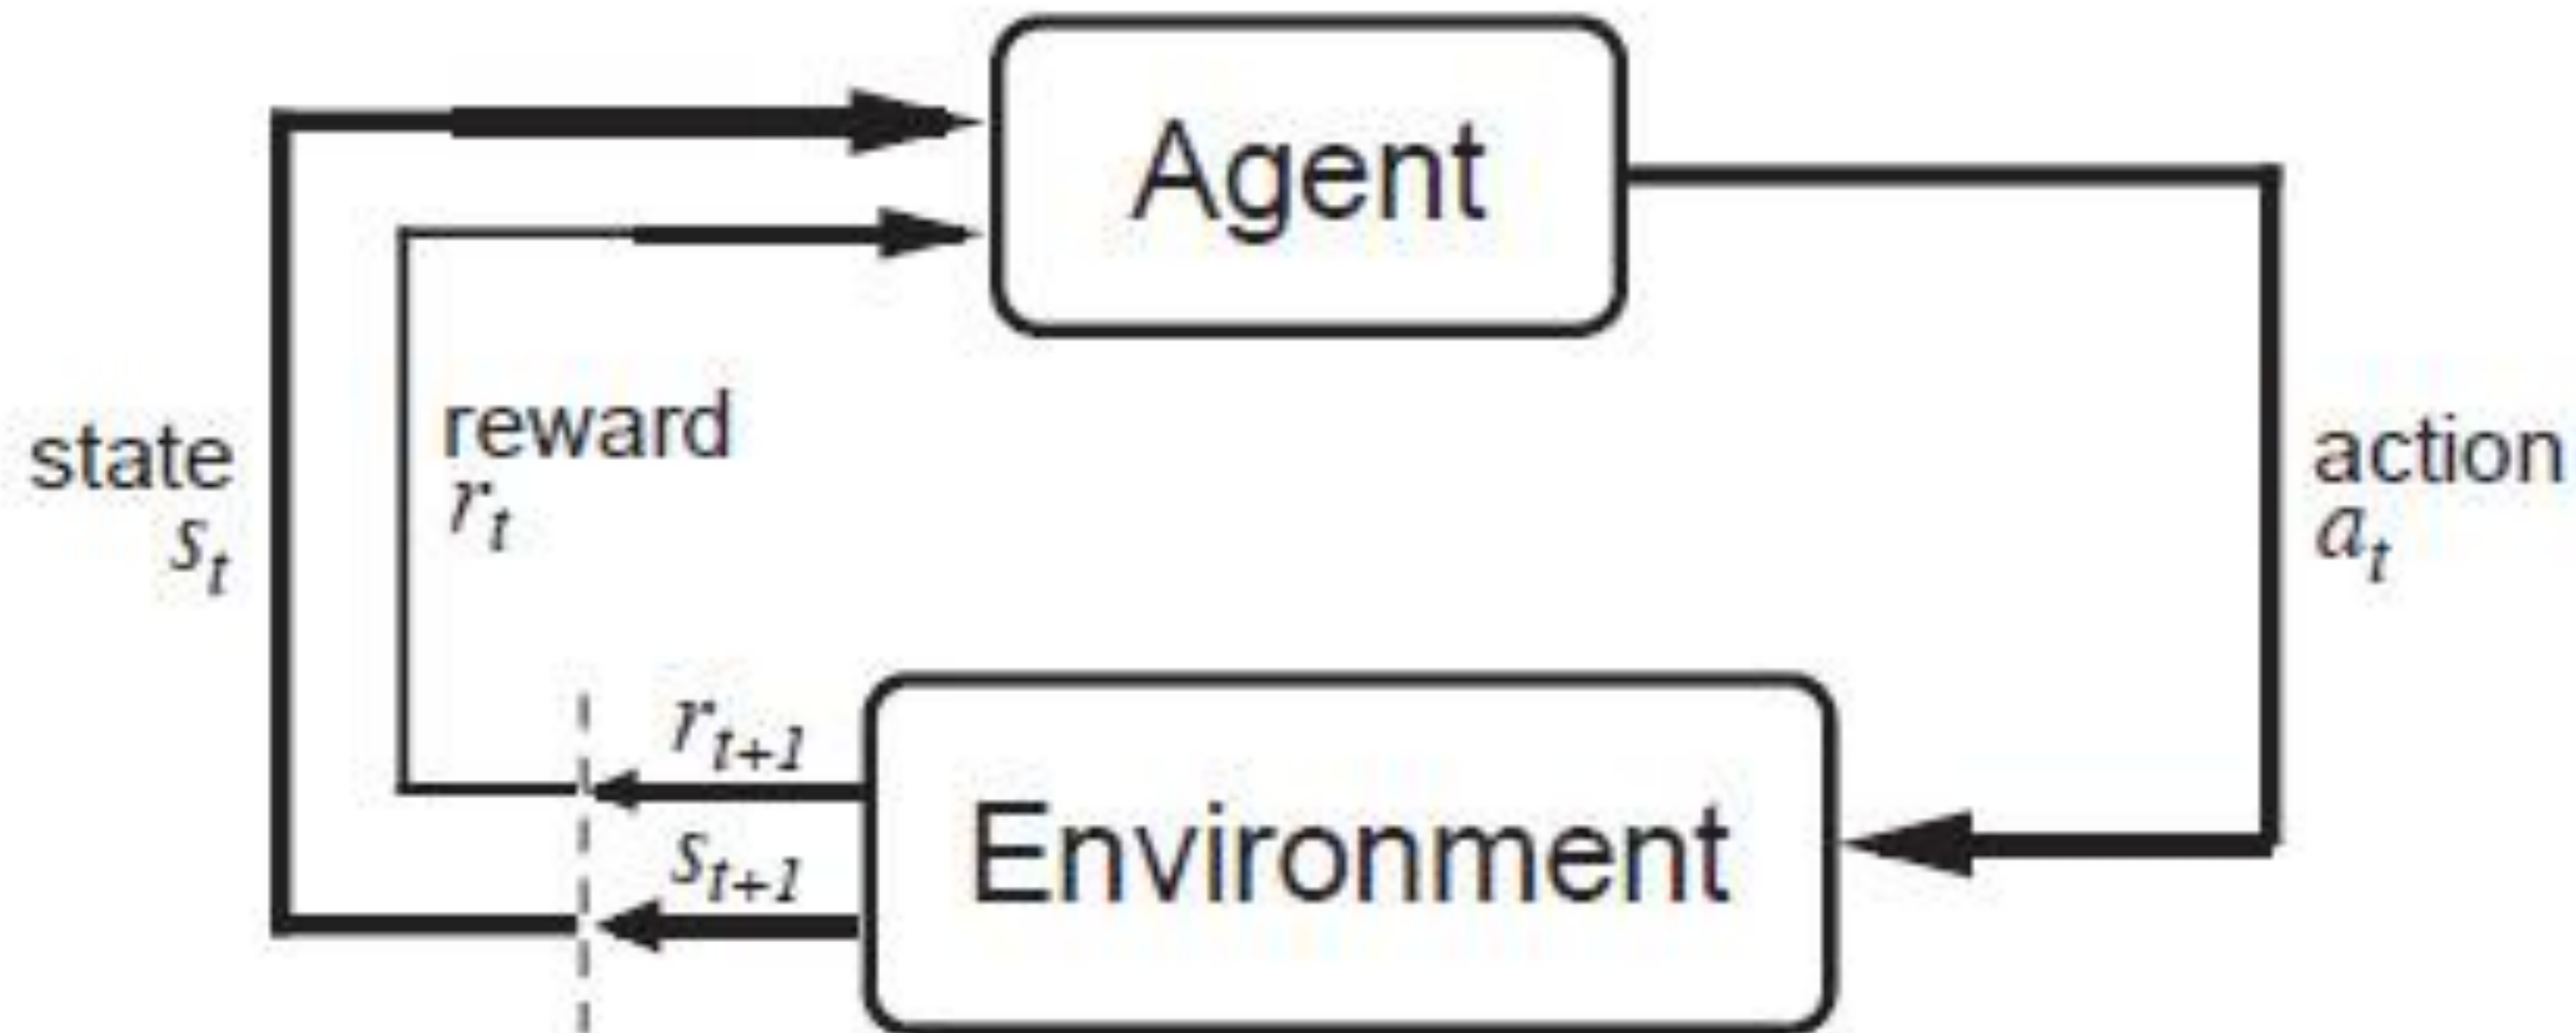
- Discrete state space : $S = \{ A , B \}$
- Discrete action space : $A = \{ X, Y \}$
- (Action conditional) State transition probability : $P(S'|S , A) = \{ \dots \}$
- Reward function : $R(S'=A) = +1$, $R(S'=B) = -1$
- Purpose : Finding an optimal policy (maximizes the expected sum of future reward)



Markov Decision Processes

- **Markov decision processes :**
a mathematical framework for modeling decision making.
- MDP are solved via [dynamic programming](#) and [reinforcement learning](#).
- Applications : [robotics](#), [automated control](#), [economics](#) and [manufacturing](#).
- Examples of MDP :
 - 1) AlphaGo에서는 바둑을 MDP로 정의
 - 2) 자동차 운전을 MDP로 정의
 - 3) 주식시장을 MDP로 정의

Agent-Environment Interaction



- Objective : Finding an optimal policy which maximizes the expected sum of future rewards
- Algorithms
 - 1) **Planning** : Exhaustive Search / Dynamic Programming
 - 2) **Reinforcement Learning** : MC method / TD Learning(Q-learning , ...)

Discount Factor

- Sum of future rewards **in episodic tasks**

➔ $G_t := R_{t+1} + R_{t+2} + R_{t+3} + \dots + R_T$

- Sum of future rewards **in continuous tasks**

➔ $G_t := R_{t+1} + R_{t+2} + R_{t+3} + \dots + R_T + \dots$
 $G_t \rightarrow \infty$ (diverge)

- Sum of discounted future rewards **in both case**


➔ $G_t := R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-1} R_T + \dots$

$$= \sum_{k=1}^{\infty} \gamma^{k-1} R_{t+k} \text{ (converge)}$$

(R_t is bounded / γ : discount factor, $0 \leq \gamma < 1$)


Deterministic Policy, $a=f(s)$

State S	$f(S)$
A	Y
B	Y




Stochastic Policy, $p(a|s)$

State S	$P(X S)$	$P(Y S)$
A	0.3	0.7
B	0.4	0.6




Gridworld



a) gridworld

3.3	8.8	4.4	5.3	1.5
1.5	3.0	2.3	1.9	0.5
0.1	0.7	0.7	0.4	-0.4
-1.0	-0.4	-0.4	-0.6	-1.2
-1.9	-1.3	-1.2	-1.4	-2.0

b) v



c) π_*

- State : 5x5
- Action : 4방향이동
- Reward : A에 도착하면 +10, B에 도착하면 +5, 벽에 부딪히면 -1, 그이외 0
- Discounted Factor : 0.9

Value-based Approach

Value Function

- We will introduce a **value function** which let us know the expected sum of future rewards at given state s , following policy π .

1) State-value function

$$v_{\pi}(s) \doteq \mathbb{E}_{\pi}[G_t \mid S_t = s] = \mathbb{E}_{\pi}\left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s\right]$$

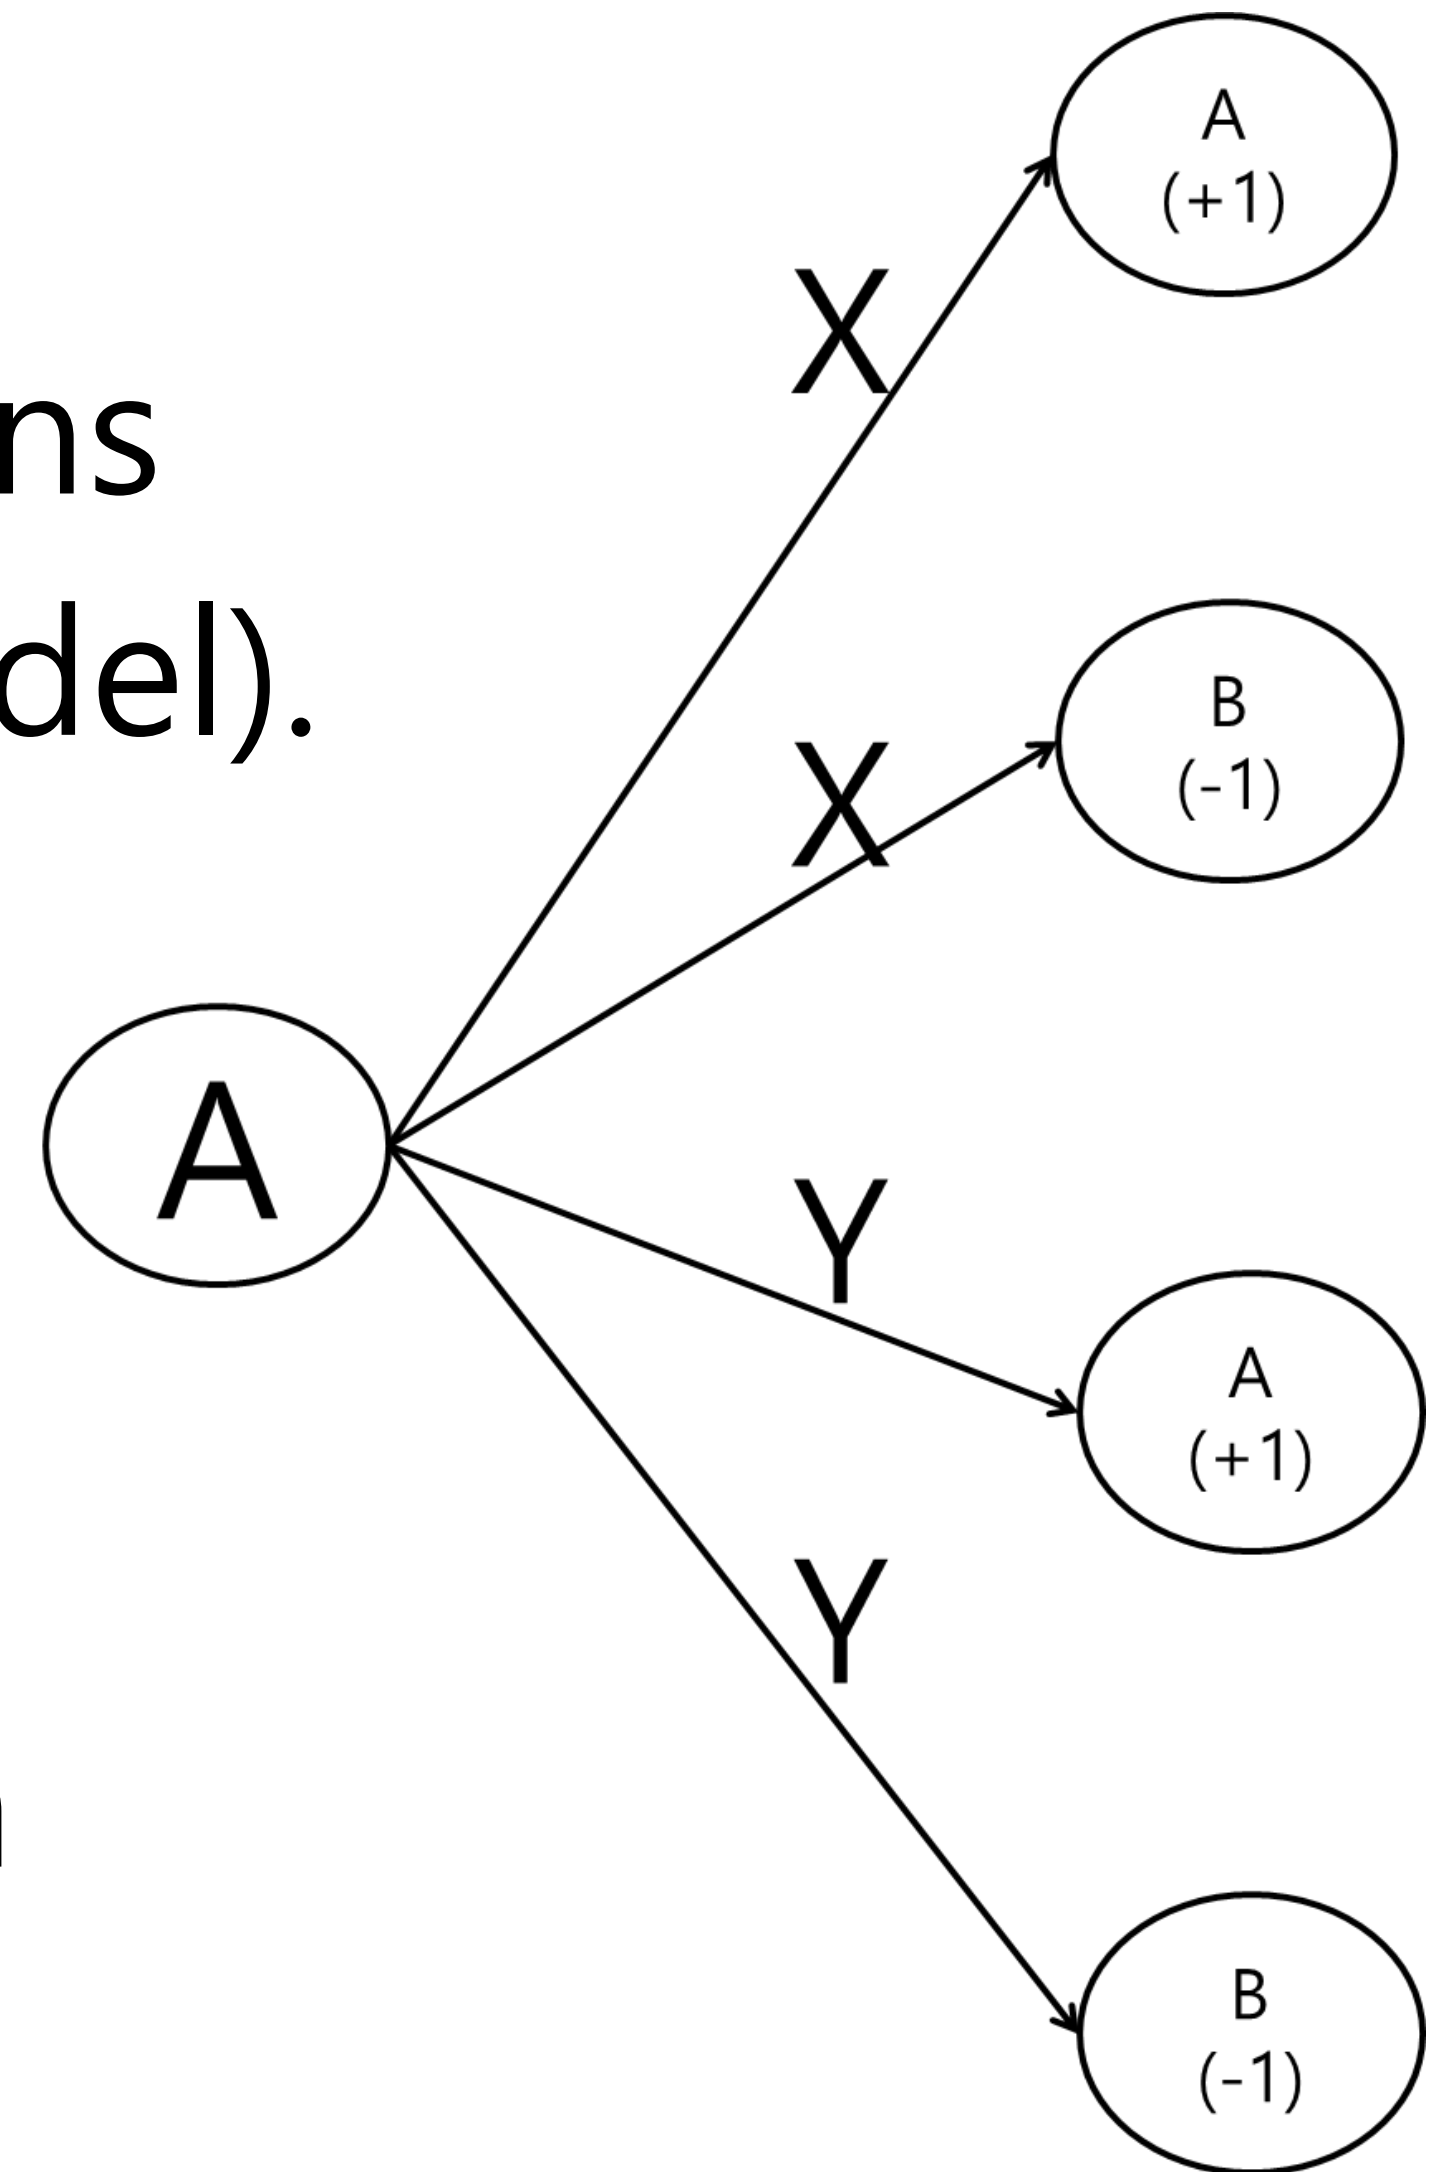
2) Action-value function

$$q_{\pi}(s, a) \doteq \mathbb{E}_{\pi}[G_t \mid S_t = s, A_t = a] = \mathbb{E}_{\pi}\left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a\right]$$

Optimal Control(Policy) with Value Function

1) Policy from state-value function

→ One-step ahead search for all actions with state transition probability(model).



2) Policy from action-value function

→ $a = f(s) = \operatorname{argmax}_a \{q_{\pi}(s, a)\}$


Planning vs Learning

- Objective : Finding an optimal policy which maximizes the expected sum of future rewards
- Algorithms
 - 1) **Planning** : Exhaustive Search / Dynamic Programming
→ 주사위 눈의 평균 = $1 \cdot 1/6 + 2 \cdot 1/6 + \dots + 6 \cdot 1/6 = 3.5$
 - 2) **Reinforcement Learning** : MC method / TD Learning
→ 주사위 눈의 평균 = 100번을 던져서 나온 눈을 평균 냄 = 3.5


Solution of the MDP : Planning

- So our last job is find optimal policy and there are two approaches.

1) Exhaustive Search





2) Dynamic Programming



Find Optimal Policy with Exhaustive Search


If we know the one step dynamics of the MDP, $P(s', r | s, a)$, we can do exhaustive search iteratively until the end step T .

And we can choose the optimal action path, but this needs $O(N^T)$!




Dynamic Programming

- 전체 문제를 풀 때 중복되는 계산이 발생 →
이를 subproblem으로 나누어 풀어 중복계산을 없앴
- 전체 path단위로 계산을 하면 앞부분 계산이 항상 중첩 →
두 step 단위로 subproblem 계산을 끝내고 다음 step으로 넘어감.



Dynamic Programming

- We can apply the DP in this problem, and the computational cost reduces to $O(N^2T)$. (But still we need to know the environment dynamics.)
- DP is a computer science technique which calculates the final goal value with compositions of cumulative partial values.



Policy Iteration

- Policy iteration consists of two simultaneous, interacting processes.
- **Policy evaluation (Iterative formulation):**
First, make the value function more precise with the current policy.
- **Policy improvement (Greedy selection):**
Second, make greedy policy greedy with respect to the current value function.

Policy Iteration

- How can we get the state-value function with DP?
(The state-value is subproblem. And action-value function is also similarly computed.)



◆ **Policy Iteration = Policy Evaluation + Policy Improvement**

$$\begin{aligned} v_{\pi}(s) &\doteq \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots \mid S_t = s] \\ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s] \\ &= \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')], \end{aligned}$$

$$\begin{aligned} \pi'(s) &\doteq \operatorname{argmax}_a q_{\pi}(s, a) \\ &= \operatorname{argmax}_a \mathbb{E}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \operatorname{argmax}_a \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')], \end{aligned}$$

Policy Iteration

$$\pi_0 \xrightarrow{\text{E}} v_{\pi_0} \xrightarrow{\text{I}} \pi_1 \xrightarrow{\text{E}} v_{\pi_1} \xrightarrow{\text{I}} \pi_2 \xrightarrow{\text{E}} \dots \xrightarrow{\text{I}} \pi_* \xrightarrow{\text{E}} v_*$$



Solution of the MDP : Learning

- The planning methods must know the perfect dynamics of environment, $P(s', r | s, a)$
- But typically this is really hard to know and empirically impossible. Therefore we will ignore this and just calculate the mean from samples.
- This is the starting point of the machine learning is embedded.

- 1) Monte Carlo Methods
- 2) Temporal-Difference Learning
(a.k.a reinforcement learning)

Monte Carlo Methods

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:

$G \leftarrow$ return following the first occurrence of s

Append G to $Returns(s)$


$V(s) \leftarrow \text{average}(Returns(s))$

Tabular state-value function

Starting State	Value
S1	Average of G(S1)
S2	Average of G(S2)
S3	Average of G(S2)

Monte Carlo Methods

- We need a full length of experience for each starting state.
- This is really time consuming to update one state while waiting the terminal of episode.
- Continuous task에는 적용 불가



Q-learning (Temporal Difference Learning)

Bellman Equation

$$\begin{aligned}v_{\pi}(s) &\doteq \mathbb{E}_{\pi}[G_t \mid S_t = s] \\&= \mathbb{E}_{\pi}\left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s\right] \\&= \mathbb{E}_{\pi}\left[R_{t+1} + \gamma \sum_{k=0}^{\infty} \gamma^k R_{t+k+2} \mid S_t = s\right] \\&= \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s] \text{ (iterative formulation)}\end{aligned}$$


$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a]$$

Temporal-Difference Learning

- TD learning is a combination of Monte Carlo idea and dynamic programming (DP) idea.
- Like Monte Carlo methods, TD methods can learn directly from raw experience without a model of the environment's dynamics.
- Like DP, TD methods update estimates based in part without waiting for a final outcome (they bootstrap).

Temporal-Difference Learning

Input: the policy π to be evaluated
Initialize $V(s)$ arbitrarily (e.g., $V(s) = 0, \forall s \in \mathcal{S}^+$)
Repeat (for each episode):
 Initialize S
 Repeat (for each step of episode):
 $A \leftarrow$ action given by π for S
 Take action A , observe R, S'
 $V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$
 $S \leftarrow S'$
 until S is terminal



Q-learning Algorithm

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right]$$

$\underbrace{Q(s_t, a_t)}_{\text{The New Action Value = The Old Value}} = \underbrace{+}_{\text{The Learning Rate}} \times \left(\underbrace{r_{t+1} + \gamma \max_a Q(s_{t+1}, a)}_{\text{The New Information}} - \underbrace{Q(s_t, a_t)}_{\text{The Old Information}} \right)$

Initialize $Q(s, a)$ arbitrarily

Repeat (for each episode):

Initialize s

Repeat (for each step of episode):

Choose a from s using policy derived from Q (e.g., ϵ -greedy)

Take action a , observe r, s'

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

$s \leftarrow s'$;

until s is terminal

Temporal-Difference Learning




Figure 6.2: Changes recommended in the driving home example by Monte Carlo methods (left) and TD methods (right).

On policy / Off policy

- **On policy** : Target policy = Behavioral policy
there can be only one policy.
→ This can learn a stochastic policy. Ex) SARSA
- **Off policy** : Target policy \neq Behavioral policy
there can be several policies.
→ Sample efficient. Ex) Q-learning

Sarsa: On-Policy TD Control

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Repeat (for each step of episode):

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

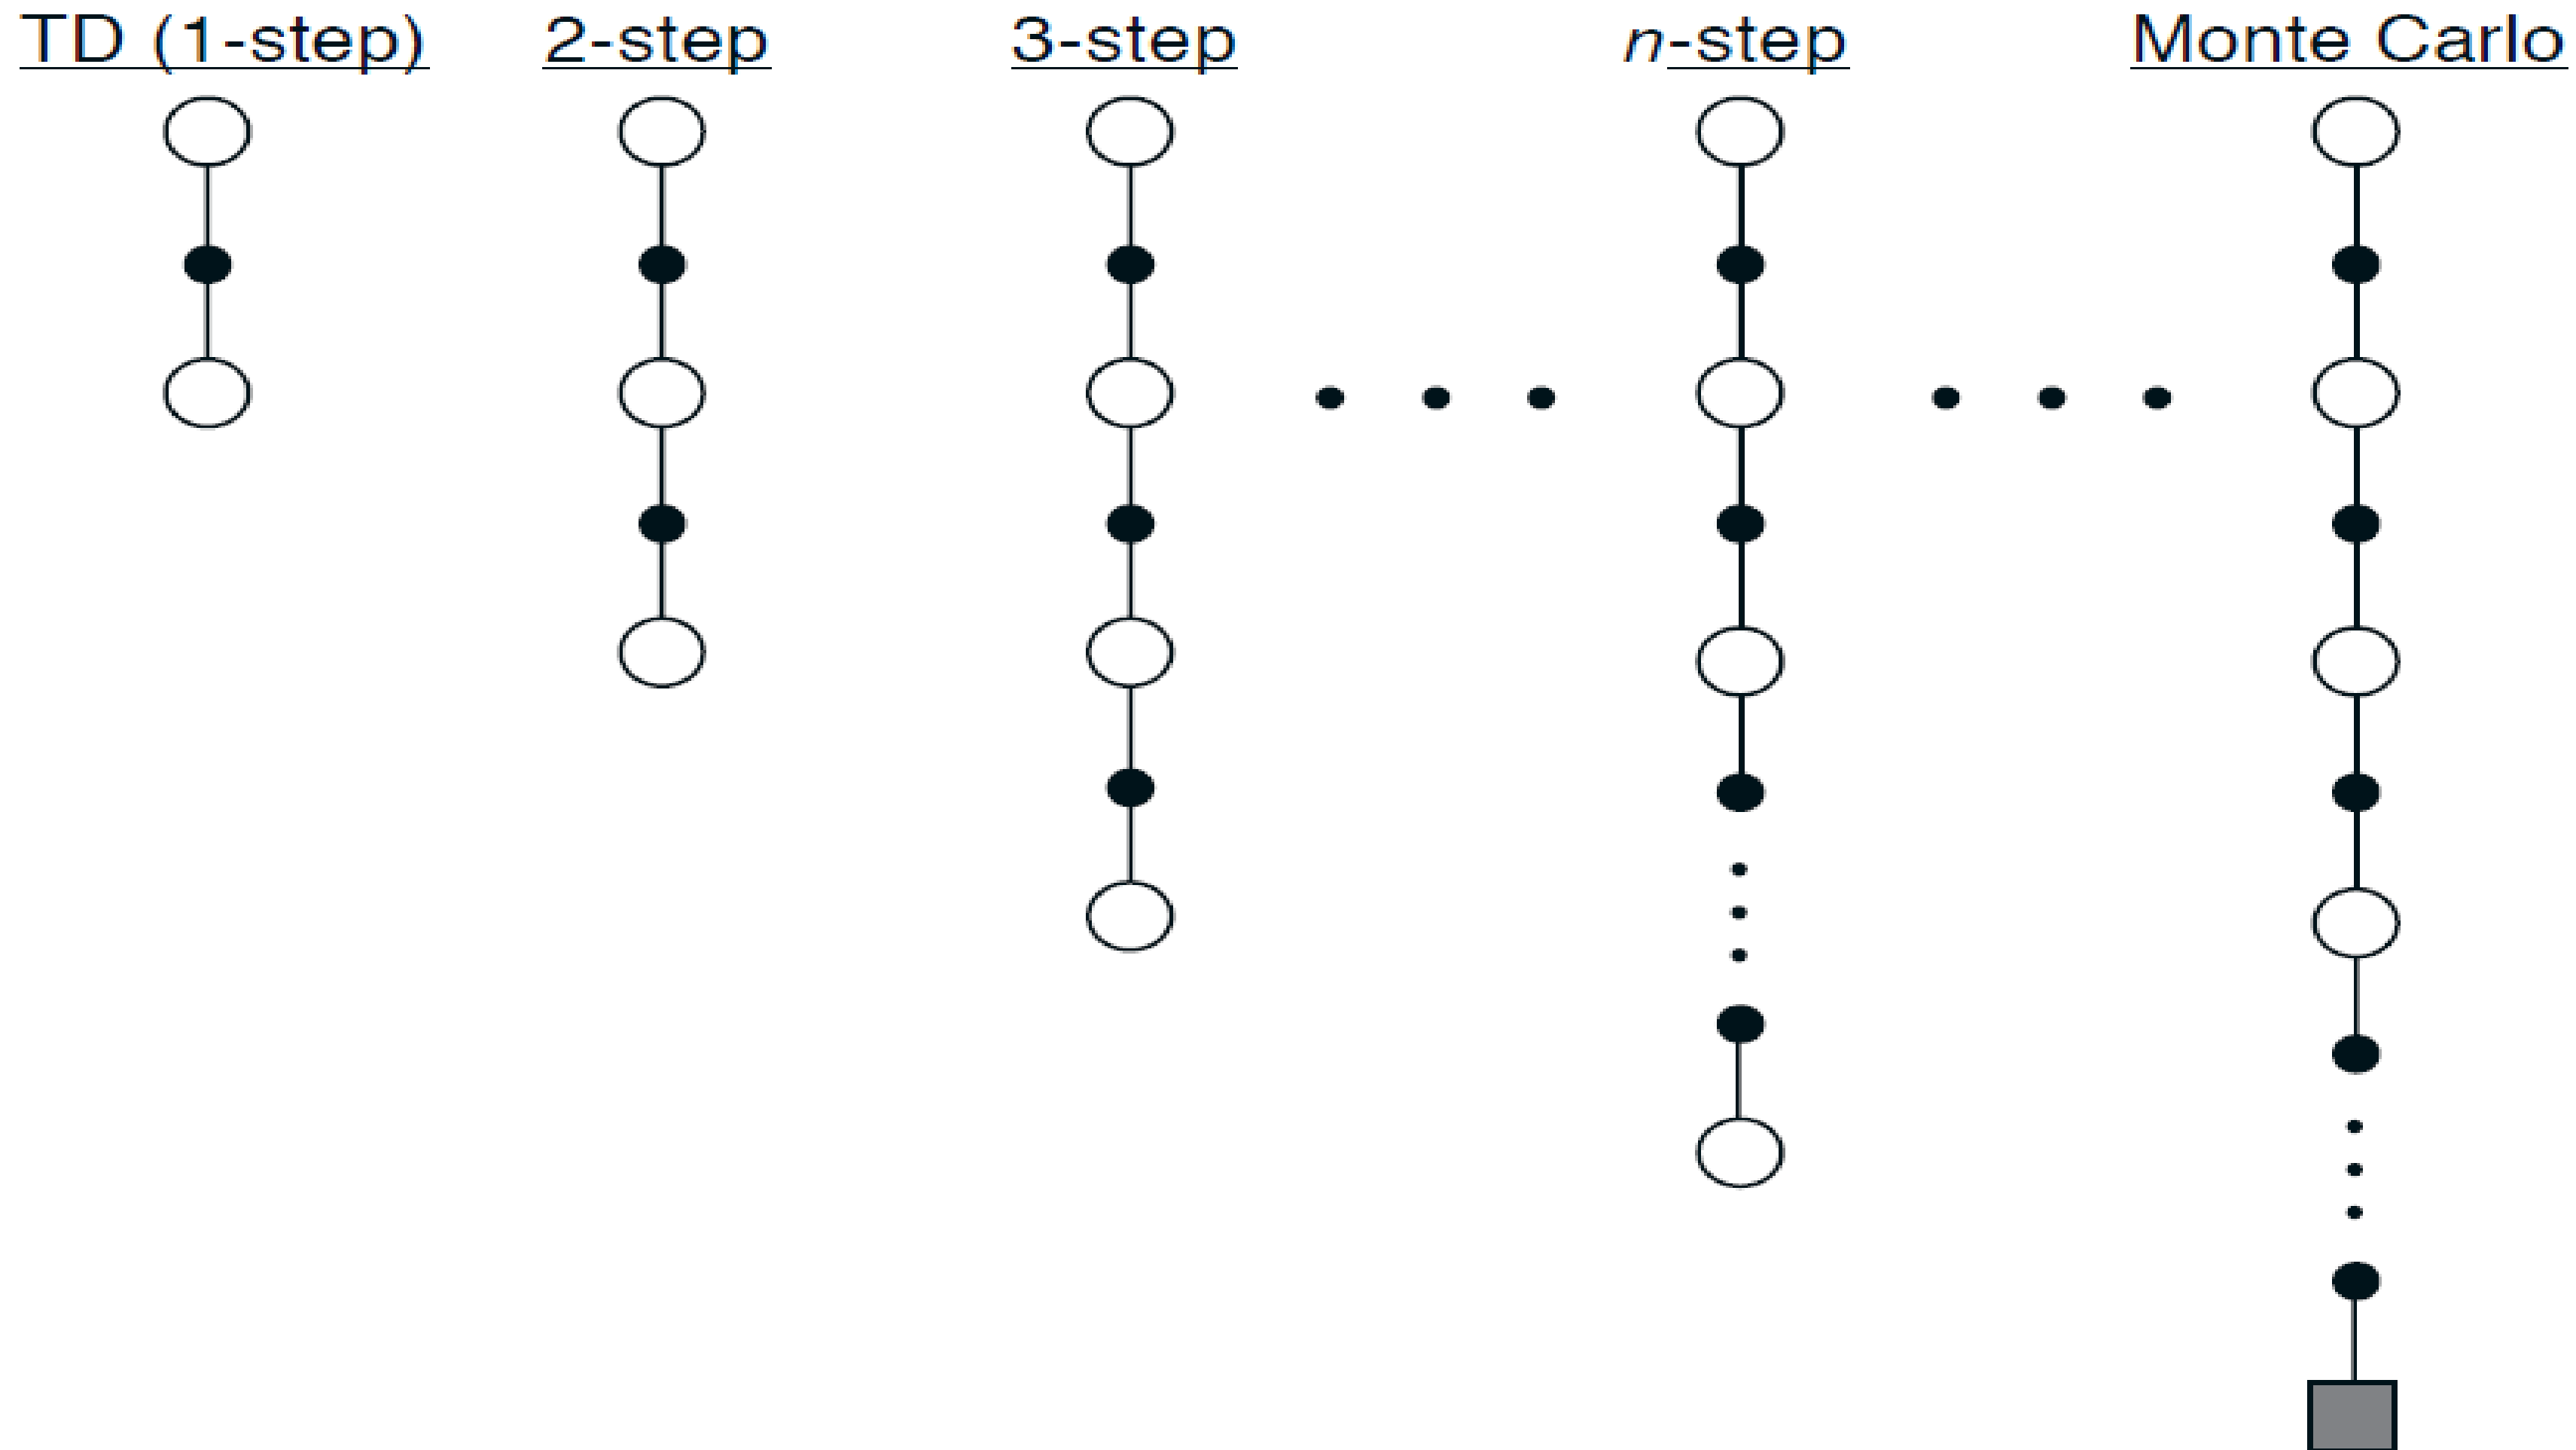
$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$


 until S is terminal

Eligibility Trace

- Smoothly combining the TD and MC.




Eligibility Trace




Comparisons


Path taken



Action values increased
by one-step Sarsa

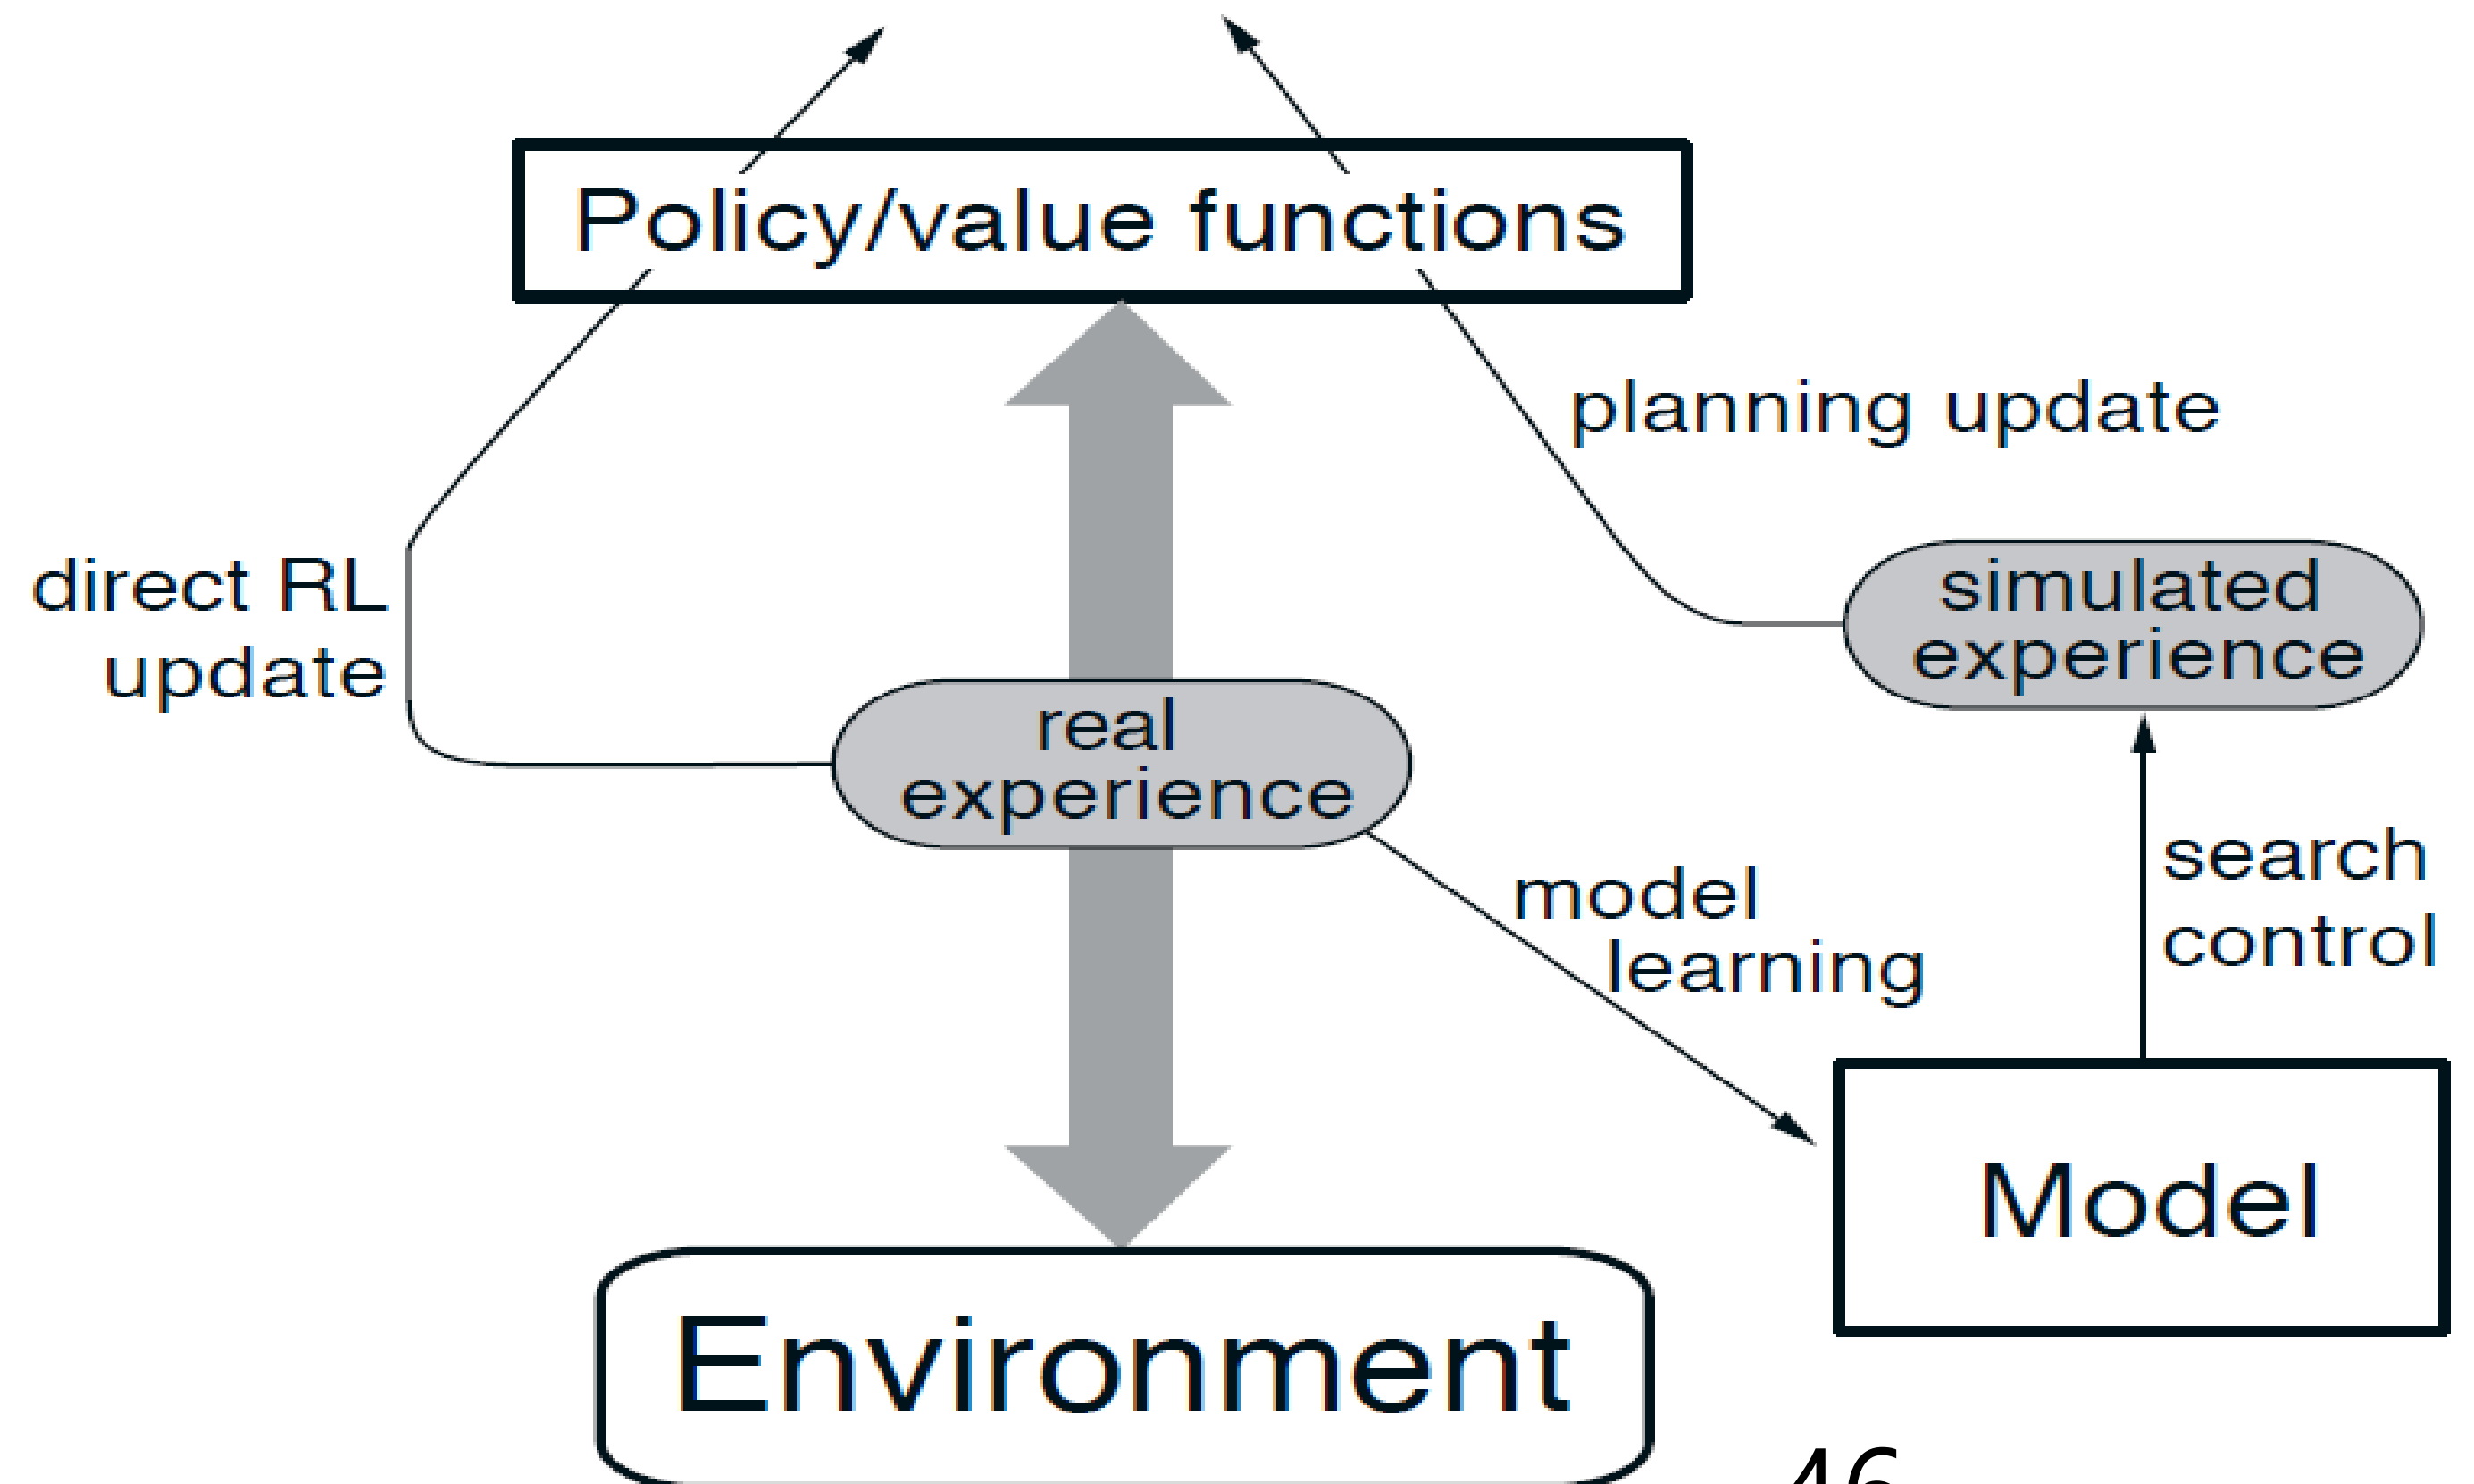



Action values increased
by Sarsa(λ) with $\lambda=0.9$



Planning + Learning

- There is only a difference between planning and learning. That is the existence of model.
- So we call planning is **model-based** method, and learning is **model-free** method.



Planning + Learning

Initialize $Q(s, a)$ and $Model(s, a)$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$
Do forever:
 (a) $S \leftarrow$ current (nonterminal) state
 (b) $A \leftarrow \epsilon$ -greedy(S, Q)
 (c) Execute action A ; observe resultant reward, R , and state, S'
 (d) $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$
 (e) $Model(S, A) \leftarrow R, S'$ (assuming deterministic environment)
 (f) Repeat n times:
 $S \leftarrow$ random previously observed state
 $A \leftarrow$ random action previously taken in S
 $R, S' \leftarrow Model(S, A)$
 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

Figure 8.4: Dyna-Q Algorithm. $Model(s, a)$ denotes the contents of the model (predicted next state and reward) for state–action pair s, a . Direct reinforcement learning, model-learning, and planning are implemented by steps (d), (e), and (f), respectively. If (e) and (f) were omitted, the remaining algorithm would be one-step tabular Q-learning.

Deep Reinforcement Learning

Generalization with deep learning

- Value function approximation with deep learning
→ **Large scale or infinite dimensional state can be solvable**


Tabular Q value function

$\hat{Q} =$

<i>state \ action</i>	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>	<i>F</i>
<i>A</i>	—	—	—	—	80	—
<i>B</i>	—	—	—	64	—	100
<i>C</i>	—	—	—	64	—	—
<i>D</i>	—	80	51	—	80	—
<i>E</i>	64	—	—	64	—	100
<i>F</i>	—	80	—	—	80	100



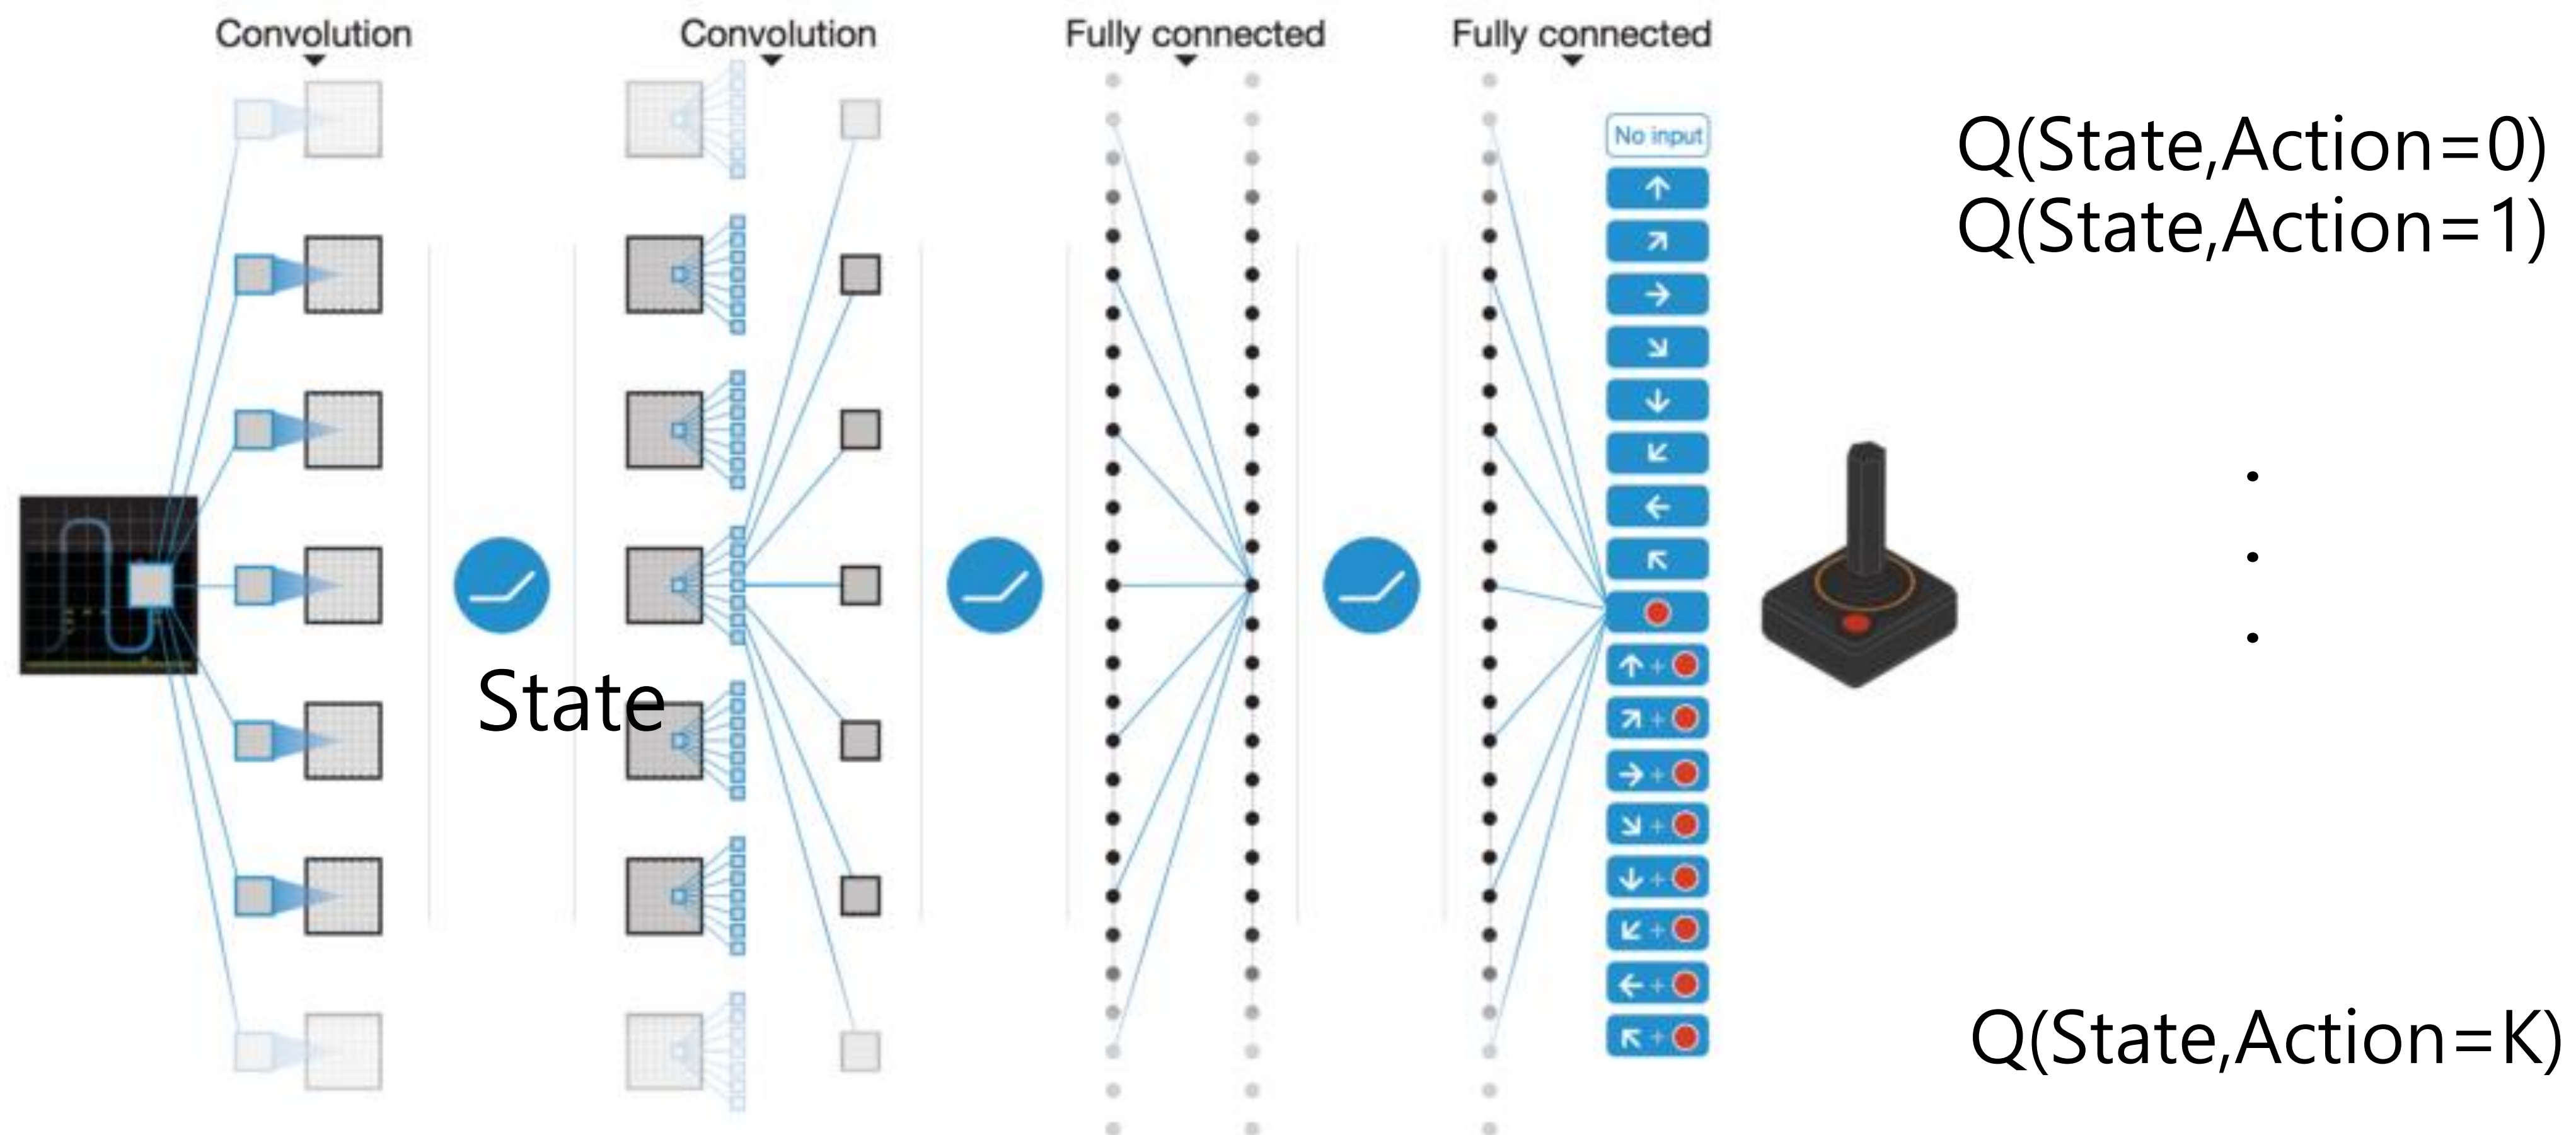
Deep neural network



This needs supervised learning techniques and online moving target regression.

Deep Q Networks (DQN)

- Q learning에서 특정 state에 대한 action-value function으로 CNN을 사용



Loss function

- Q learning의 업데이트 식

TD error

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right]$$

- TD error의 크기를 Loss로 사용, regression 방식으로 학습

$$L = \frac{1}{2} \left[\underbrace{r + \max_{a'} Q(s', a')}_{\text{target}} - \underbrace{Q(s, a)}_{\text{prediction}} \right]^2$$

Target Network

$$L = \frac{1}{2} \left[\underbrace{r + \max_{a'} Q(s', a')}_{\text{target}} - \underbrace{Q(s, a)}_{\text{prediction}} \right]^2$$

- **Moving Target Regression :**

Q를 학습하는 순간, target 값도 같이 변화 → 학습 성능 저하

➔ **Target Network** : DQN과 동일한 구조를 가지고 있으며 학습 도중 weight값이 변하지 않는 별도의 네트워크로 $Q(s', a')$ 로 **고정된 target value**를 사용하여 학습 안정성 개선

(Target Network의 weight값들은 주기적으로 DQN의 것을 복사)

Replay Memory

- **Data inefficiency** : 그 동안 경험한 experience들을 off-policy learning으로 재활용(Mini-batch learning)
- **Data correlation** : Mini-batch를 On-line으로 구성할 경우 Mini batch 내의 데이터들이 서로 비슷함(Correlated)
→ 한쪽으로 치우친 불균형한 학습이 발생함
- **Replay Memory** : (State, Action, Reward, Next State) 데이터를 Buffer에 저장해놓고 그 안에서 Random Sampling하여 Mini-batch를 구성하는 방법으로 해결

Reward Clipping


- **Reward Scale problem** : 도메인에 따라 Reward의 scale이 다르기 때문에 학습되는 Q-value의 크기도 제각각.

이렇게 학습해야하는 Q-value의 variance가 매우 큰 경우
(ex : -1000~1000) 신경망 학습이 어려움.


➔ **Reward Clipping**: Reward의 크기를 -1~1 사이로 잘라내어 안정적으로 학습

DQN의 성능

- ATARI 2600 고전게임에서 실험
- 절반 이상의 게임에서 사람보다 우수
- 기존방식 (linear)에 비해 월등한 향상
- 일부 게임은 학습에 실패함.
(sparse reward, 복잡한 단계를 가진 게임)



DQN 데모



DQN의 발전된 모델들 - Double DQN

$$Y_t^{\text{DQN}} \equiv R_{t+1} + \gamma \max_a Q(S_{t+1}, a; \theta_t^-).$$

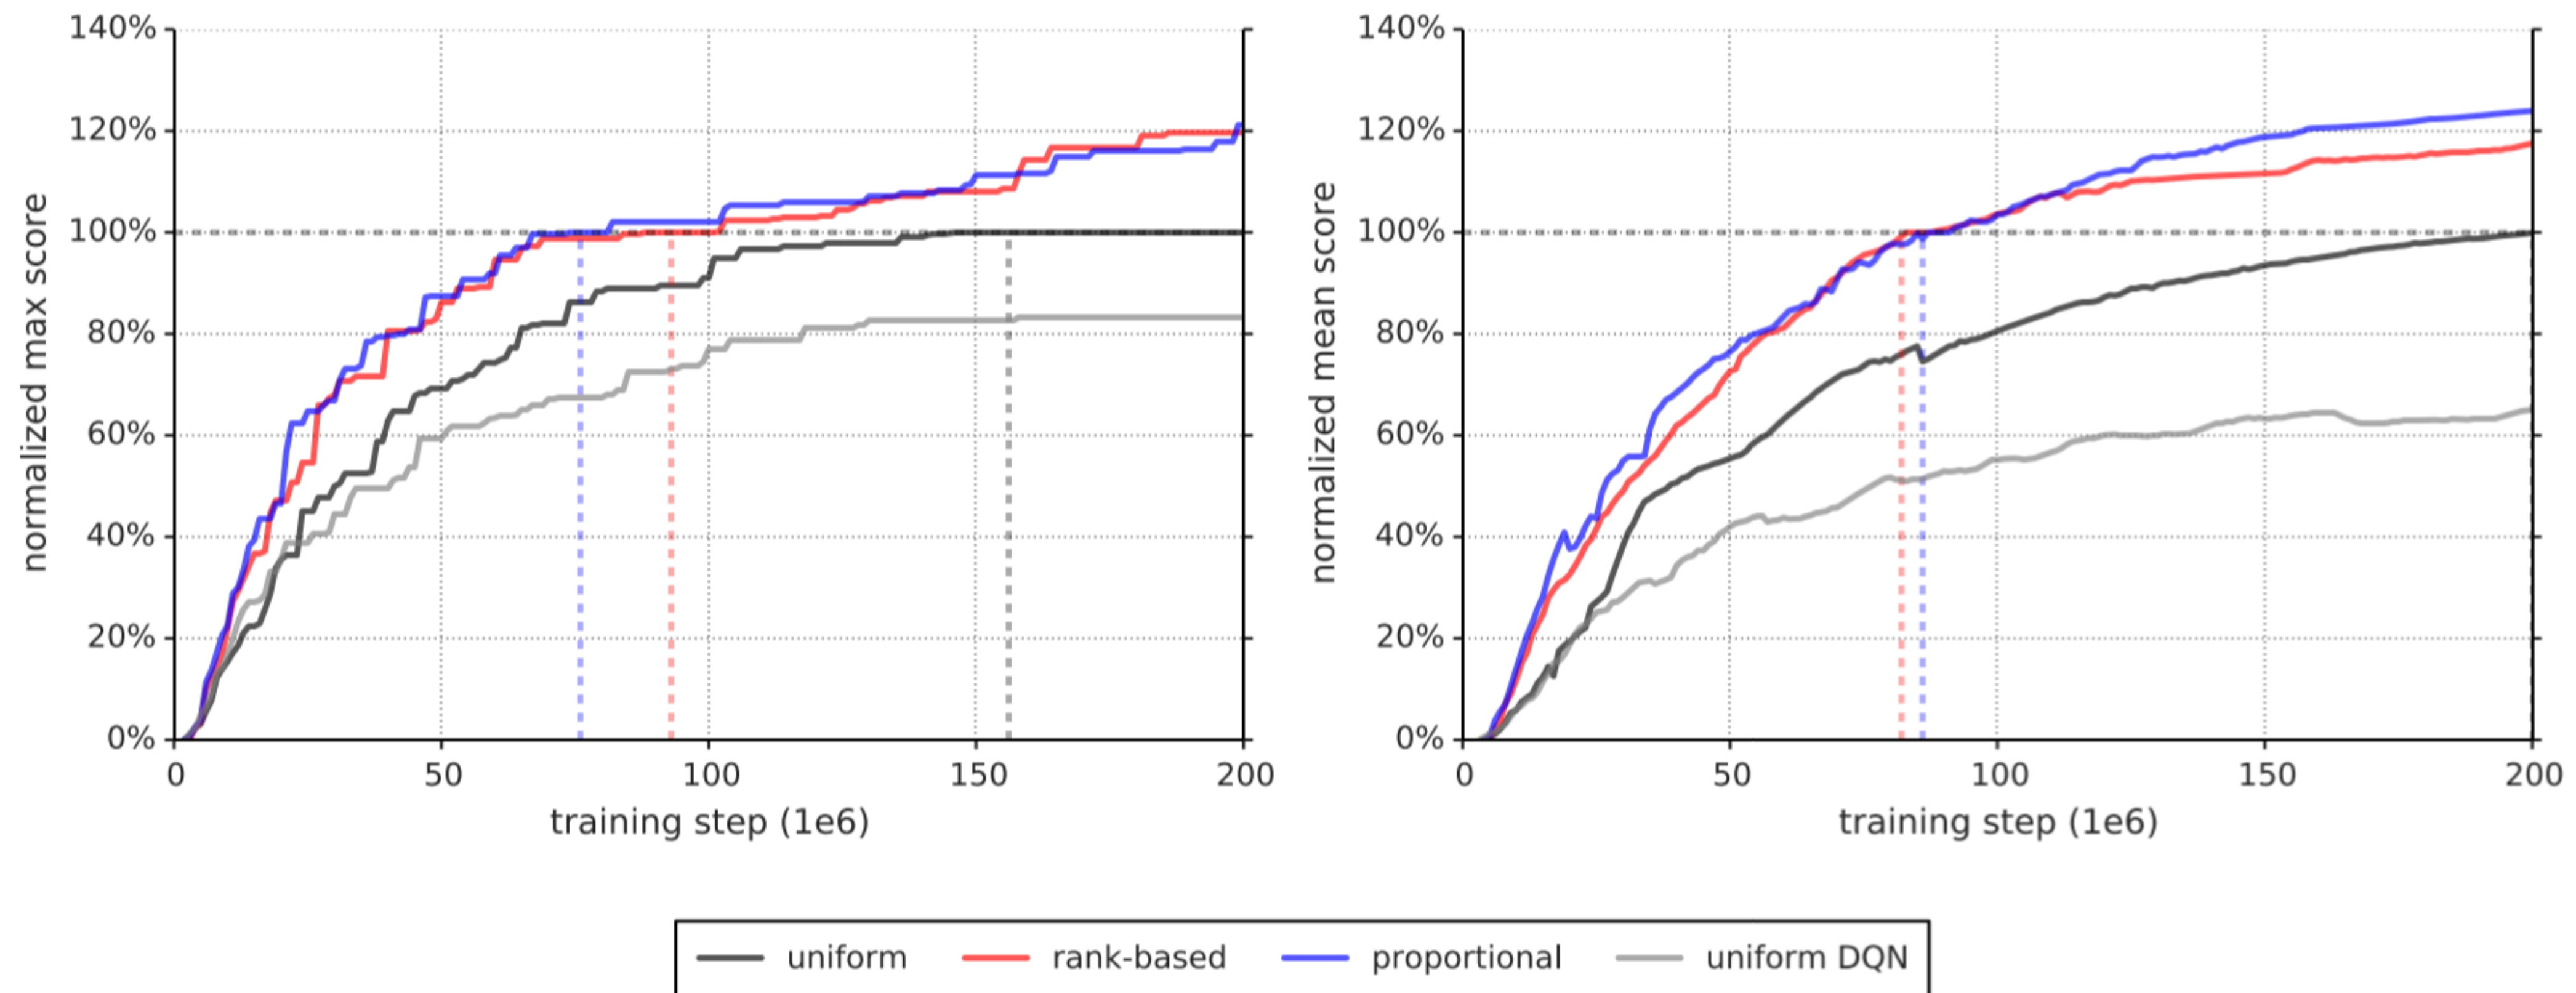
$$Y_t^{\text{DoubleQ}} \equiv R_{t+1} + \gamma Q(S_{t+1}, \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a; \theta_t); \theta'_t).$$

• **Positive bias:** max 연산자에 의해 Q-value를 실제보다 높게 평가하는 경향이 있음

➔ Double Q learning을 사용하여 이 현상을 방지, DQN에 보다 빠른 학습이 가능

DQN의 발전된 모델들 - Prioritized Replay Memory

- 일반적인 DQN은 Replay Memory에서 Uniform Sampling
➔ TD error가 높은 데이터들이 선택될 확률을 높여,
더 빠르고 효율적인 학습이 가능



References

- [1] Sutton, Richard S., and Andrew G. Barto. Reinforcement learning: An introduction. MIT press, 1998.
- [2] Mnih, Volodymyr, et al. "Human-level control through deep reinforcement learning." Nature 518.7540 (2015): 529-533.
- [3] Van Hasselt, Hado, Arthur Guez, and David Silver. "Deep Reinforcement Learning with Double Q-Learning." AAAI. 2016.
- [4] Wang, Ziyu, et al. "Dueling network architectures for deep reinforcement learning." arXiv preprint arXiv:1511.06581 (2015).
- [5] Schaul, Tom, et al. "Prioritized experience replay." arXiv preprint arXiv:1511.05952 (2015).

Appendix

- Atari 2600 - <https://www.youtube.com/watch?v=iqXKQf2BOSE>
- Super MARIO - <https://www.youtube.com/watch?v=qv6UVOQ0F44>
- Robot Learns to Flip Pancakes - https://www.youtube.com/watch?v=W_gxLKSsSIE
- Stanford Autonomous Helicopter - Airshow #2 - <https://www.youtube.com/watch?v=VCdxqn0fcnE>

- OpenAI Gym - <https://gym.openai.com/envs>
- Awesome RL - <https://github.com/aikorea/awesome-rl>
- Udacity RL course
- TensorFlow DRL - <https://github.com/nivwusquorum/tensorflow-deepq>
- Karpathy rldemo - <http://cs.stanford.edu/people/karpathy/convnetjs/demo/rldemo.html>

감사합니다.